

Import the necessary libraries

```
import tensorflow as tf
import numpy as np
import os
from tensorflow.keras.preprocessing.image import
ImageDataGenerator,load_img
import matplotlib.pyplot as plt
from keras.models import Sequential,Model
from keras.layers import
Dense,Flatten,Conv2D,MaxPooling2D,Dropout,BatchNormalization
from keras.callbacks import EarlyStopping
from sklearn.metrics import confusion_matrix,classification_report
from keras.applications import resnet_v2
import keras
import seaborn as sns

from google.colab import drive
drive.mount('/content/drive')

Mounted at /content/drive

from zipfile import ZipFile
file_name = "/content/drive/MyDrive/Colab Files/data.zip"

with ZipFile(file_name,'r') as zip:
    zip.extractall()
    print('Done')

Done
```

Plot the sample images for all the classes

Plot the distribution of images across the classes

```
picture_size = 48
folder_path = '/content/data/'

expression = 'happy'

plt.figure(figsize= (10,10))
for i in range(1, 10, 1):
    plt.subplot(3,3,i)
```

```
img = load_img(folder_path+"train/"+expression+"/"+  
os.listdir(folder_path + "train/" + expression)[i],  
               target_size=(picture_size, picture_size))  
plt.imshow(img)  
plt.axis('off')  
plt.show()
```



```
train_dir = '/content/data/train'  
test_dir = '/content/data/test'  
  
train_angry_dir = '/content/data/train/angry'  
train_disgust_dir = '/content/data/train/disgust'
```

```
train_fear_dir = '/content/data/train/fear'
train_happy_dir = '/content/data/train/happy'
train_neutral_dir = '/content/data/train/neutral'
train_sad_dir = '/content/data/train/sad'
train_surprise_dir = '/content/data/train/surprise'

test_angry_dir = '/content/data/test/angry'
test_disgust_dir = '/content/data/test/disgust'
test_fear_dir = '/content/data/test/fear'
test_happy_dir = '/content/data/test/happy'
test_neutral_dir = '/content/data/test/neutral'
test_sad_dir = '/content/data/test/sad'
test_surprise_dir = '/content/data/test/surprise'

dir_list =
[train_angry_dir,train_disgust_dir,train_fear_dir,train_happy_dir,train_neutral_dir,train_sad_dir,

train_surprise_dir,test_angry_dir,test_disgust_dir,test_fear_dir,test_happy_dir,test_neutral_dir,
    test_sad_dir,test_surprise_dir]

for d in dir_list:
    print(d,len(os.listdir(d)))

/content/data/train/angry 3992
/content/data/train/disgust 436
/content/data/train/fear 4103
/content/data/train/happy 7164
/content/data/train/neutral 4982
/content/data/train/sad 4938
/content/data/train/surprise 3205
/content/data/test/angry 960
/content/data/test/disgust 111
/content/data/test/fear 1018
/content/data/test/happy 1825
/content/data/test/neutral 1216
/content/data/test/sad 1139
/content/data/test/surprise 797
```

Build a data augmentation for train data to create new data with translation, rescale and flip, and rotation transformations. Rescale the image at 48x48

Build a data augmentation for test data to create new data and rescale the image at 48x48

Read images directly from the train folder and test folder using the appropriate function

```
# image generator
train_datagen = ImageDataGenerator(rescale=1.0/255.0,
                                   rotation_range=40,
                                   width_shift_range=0.2,
                                   height_shift_range=0.2,
                                   zoom_range=0.2,
                                   horizontal_flip=True,
                                   fill_mode='nearest')

train_generator = train_datagen.flow_from_directory(train_dir,
target_size=(150,150),
                                                    batch_size=64,

class_mode='categorical')

validation_datagen = ImageDataGenerator(rescale=1.0/255.0)

validation_generator =
validation_datagen.flow_from_directory(test_dir,

target_size=(150,150),

batch_size=62,

class_mode='categorical')

Found 28820 images belonging to 7 classes.
Found 7066 images belonging to 7 classes.
```

1. CNN Architecture:

- Add convolutional layers, max pool layers, dropout layers, batch normalization layers
- Use Relu as activation functions
- Take loss function as categorical cross-entropy
- Take Adam as an optimizer
- Use early-stop with two patients and monitor for validation loss
- Try with ten number epochs
- Train the model using the generator and test the accuracy of the test data at every epoch
- Plot the training and validation accuracy, and the loss
- Observe the precision, recall the F1-score for all classes for both grayscale and color models, and determine if the model's classes are good

```
model_adam = Sequential()

# 1st CNN layer
model_adam.add(Conv2D(16,
(3,3),activation='relu',input_shape=(150,150,3)))
model_adam.add(BatchNormalization())
model_adam.add(MaxPooling2D(2,2))
model_adam.add(Dropout(0.25))

# 2nd CNN Layer
model_adam.add(Conv2D(32,(3,3),activation='relu'))
model_adam.add(BatchNormalization())
model_adam.add(MaxPooling2D(2,2))
model_adam.add(Dropout(0.25))

# 3rd CNN layer
model_adam.add(Conv2D(64,(3,3),activation='relu'))
model_adam.add(BatchNormalization())
model_adam.add(MaxPooling2D(2,2))
model_adam.add(Dropout(0.25))

# 4th CNN Layer
model_adam.add(Conv2D(64,(3,3),activation='relu'))
model_adam.add(BatchNormalization())
model_adam.add(MaxPooling2D(2,2))
model_adam.add(Dropout(0.25))
```

```

model_adam.add(Flatten())

# Fully connected layer
model_adam.add(Dense(512,activation='relu'))

model_adam.add(Dense(7,activation='softmax'))

model_adam.compile(optimizer=keras.optimizers.adam_v2.Adam(learning_rate=0.001),
                    loss=keras.losses.categorical_crossentropy,
                    metrics=['accuracy'])

```

```
model_adam.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 148, 148, 16)	448
batch_normalization (Batch Normalization)	(None, 148, 148, 16)	64
max_pooling2d (MaxPooling2D)	(None, 74, 74, 16)	0
dropout (Dropout)	(None, 74, 74, 16)	0
conv2d_1 (Conv2D)	(None, 72, 72, 32)	4640
batch_normalization_1 (Batch Normalization)	(None, 72, 72, 32)	128
max_pooling2d_1 (MaxPooling2D)	(None, 36, 36, 32)	0
dropout_1 (Dropout)	(None, 36, 36, 32)	0
conv2d_2 (Conv2D)	(None, 34, 34, 64)	18496
batch_normalization_2 (Batch Normalization)	(None, 34, 34, 64)	256
max_pooling2d_2 (MaxPooling2D)	(None, 17, 17, 64)	0
dropout_2 (Dropout)	(None, 17, 17, 64)	0
conv2d_3 (Conv2D)	(None, 15, 15, 64)	36928

batch_normalization_3 (Batch Normalization)	(None, 15, 15, 64)	256
max_pooling2d_3 (MaxPooling2D)	(None, 7, 7, 64)	0
dropout_3 (Dropout)	(None, 7, 7, 64)	0
flatten (Flatten)	(None, 3136)	0
dense (Dense)	(None, 512)	1606144
dense_1 (Dense)	(None, 7)	3591

```

=====
Total params: 1,670,951
Trainable params: 1,670,599
Non-trainable params: 352

```

```

early_stopping =
EarlyStopping(patience=2,monitor='val_loss',restore_best_weights=True)

```

```

# Fit the model

```

```

history_adam =
model_adam.fit(train_generator,epochs=10,batch_size=512,verbose=1,validation_data=validation_generator,
               callbacks=early_stopping)

```

```

Epoch 1/10

```

```

451/451 [=====] - 155s 322ms/step - loss: 1.9825 - accuracy: 0.2357 - val_loss: 1.8593 - val_accuracy: 0.2632

```

```

Epoch 2/10

```

```

451/451 [=====] - 143s 317ms/step - loss: 1.7789 - accuracy: 0.2579 - val_loss: 1.7534 - val_accuracy: 0.2714

```

```

Epoch 3/10

```

```

451/451 [=====] - 143s 317ms/step - loss: 1.7552 - accuracy: 0.2730 - val_loss: 2.0191 - val_accuracy: 0.2829

```

```

Epoch 4/10

```

```

451/451 [=====] - 144s 318ms/step - loss: 1.7077 - accuracy: 0.3068 - val_loss: 1.7208 - val_accuracy: 0.3006

```

```

Epoch 5/10

```

```

451/451 [=====] - 143s 318ms/step - loss: 1.6553 - accuracy: 0.3347 - val_loss: 1.5006 - val_accuracy: 0.4181

```

```

Epoch 6/10

```

```

451/451 [=====] - 142s 314ms/step - loss: 1.6104 - accuracy: 0.3633 - val_loss: 1.6171 - val_accuracy: 0.3913

```

```

Epoch 7/10

```

```

451/451 [=====] - 142s 315ms/step - loss: 1.5699 - accuracy: 0.3817 - val_loss: 1.5449 - val_accuracy: 0.4014

```

```

# Evaluate the model
test_loss, test_acc =
model_adam.evaluate(validation_generator,verbose=1)


print('Model Accuracy',test_acc)
print('Model Loss',test_loss)

114/114 [=====] - 5s 41ms/step - loss: 1.5006
- accuracy: 0.4181
Model Accuracy 0.418058305978775
Model Loss 1.5005847215652466

# Plot the loss function for the model
plt.plot(history_adam.history['loss'], label='train')
plt.plot(history_adam.history['val_loss'], label='test')

plt.title('Model loss: Adam optimizer, Relu Activation')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='best')
plt.show()

```



```


# Plot the accuracy function for the model
plt.plot(history_adam.history['accuracy'], label='train')
plt.plot(history_adam.history['val_accuracy'], label='test')

plt.title('Model Accuracy: Adam optimizer, Relu Activation')

```



```
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='best')
plt.show()
```



Customized CNN Architecture:

- Add convolutional layers, max pool layers, dropout layers, batch normalization layers on the top of the first model architecture to improve the accuracy
- Change the batch size activation function and optimizer as rmsprop and observe if the accuracy increases
- Take the loss function as categorical cross-entropy
- Use early stopping with the patience of two epochs and monitoring of validation loss
- Try with ten number epochs
- Train the model using the generator and test the accuracy of the test data at every epoch
- Plot the training and validation accuracy, and the loss
- Observe the precision, recall the F1-score for all classes for both grayscale and color models, and determine if the model's classes are good

```

model_rmsprop = Sequential()

# 1st CNN layer
model_rmsprop.add(Conv2D(16,
(3,3),activation='relu',input_shape=(150,150,3)))
model_rmsprop.add(BatchNormalization())
model_rmsprop.add(MaxPooling2D(2,2))
model_rmsprop.add(Dropout(0.25))

# 2nd CNN Layer
model_rmsprop.add(Conv2D(32,(3,3),activation='relu'))
model_rmsprop.add(BatchNormalization())
model_rmsprop.add(MaxPooling2D(2,2))
model_rmsprop.add(Dropout(0.25))

# 3rd CNN layer
model_rmsprop.add(Conv2D(64,(3,3),activation='relu'))
model_rmsprop.add(BatchNormalization())
model_rmsprop.add(MaxPooling2D(2,2))
model_rmsprop.add(Dropout(0.25))

model_rmsprop.add(Flatten())

# Fully connected layer
model_rmsprop.add(Dense(512,activation='relu'))

model_rmsprop.add(Dense(7,activation='softmax'))

model_rmsprop.compile(optimizer=keras.optimizers.rmsprop_v2.RMSProp(learning_rate=0.0001),
                        loss=keras.losses.categorical_crossentropy,
                        metrics=['accuracy'])

model_rmsprop.summary()

```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====		
conv2d_4 (Conv2D)	(None, 148, 148, 16)	448
batch_normalization_4 (Batch Normalization)	(None, 148, 148, 16)	64
max_pooling2d_4 (MaxPooling2D)	(None, 74, 74, 16)	0
dropout_4 (Dropout)	(None, 74, 74, 16)	0
conv2d_5 (Conv2D)	(None, 72, 72, 32)	4640

batch_normalization_5 (Batch Normalization)	(None, 72, 72, 32)	128
max_pooling2d_5 (MaxPooling2D)	(None, 36, 36, 32)	0
dropout_5 (Dropout)	(None, 36, 36, 32)	0
conv2d_6 (Conv2D)	(None, 34, 34, 64)	18496
batch_normalization_6 (Batch Normalization)	(None, 34, 34, 64)	256
max_pooling2d_6 (MaxPooling2D)	(None, 17, 17, 64)	0
dropout_6 (Dropout)	(None, 17, 17, 64)	0
flatten_1 (Flatten)	(None, 18496)	0
dense_2 (Dense)	(None, 512)	9470464
dense_3 (Dense)	(None, 7)	3591

```

=====
Total params: 9,498,087
Trainable params: 9,497,863
Non-trainable params: 224

```

```

early_stopping =
EarlyStopping(patience=2,monitor='val_loss',restore_best_weights=True)

```

```

# Fit the model

```

```

history_rmsprop =
model_rmsprop.fit(train_generator,epochs=10,batch_size=512,verbose=1,validation_data=validation_generator,
                  callbacks=early_stopping)

```

```

Epoch 1/10

```

```

451/451 [=====] - 145s 318ms/step - loss: 1.9207 - accuracy: 0.2339 - val_loss: 3.0168 - val_accuracy: 0.2349

```

```

Epoch 2/10

```

```

451/451 [=====] - 144s 320ms/step - loss: 1.7985 - accuracy: 0.2535 - val_loss: 1.9633 - val_accuracy: 0.2772

```

```

Epoch 3/10

```

```

451/451 [=====] - 142s 314ms/step - loss: 1.7827 - accuracy: 0.2635 - val_loss: 2.0487 - val_accuracy: 0.2880

```

```

Epoch 4/10

```

```

451/451 [=====] - 145s 321ms/step - loss:

```

```

1.7697 - accuracy: 0.2734 - val_loss: 1.9202 - val_accuracy: 0.3147
Epoch 5/10
451/451 [=====] - 143s 317ms/step - loss:
1.7615 - accuracy: 0.2812 - val_loss: 2.2899 - val_accuracy: 0.3135
Epoch 6/10
451/451 [=====] - 143s 316ms/step - loss:
1.7502 - accuracy: 0.2865 - val_loss: 3.3217 - val_accuracy: 0.2833

# Evaluate the model
test_loss, test_acc =
model_rmsprop.evaluate(validation_generator, verbose=1)


print('Model Accuracy', test_acc)
print('Model Loss', test_loss)

114/114 [=====] - 5s 40ms/step - loss: 1.9202
- accuracy: 0.3147
Model Accuracy 0.3147466778755188
Model Loss 1.9201639890670776

# Plot the loss function for the model
plt.plot(history_rmsprop.history['loss'], label='train')
plt.plot(history_rmsprop.history['val_loss'], label='test')


plt.title('Model loss: RMSProp optimizer, Relu Activation')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='best')
plt.show()

```



```
# Plot the accuracy function for the model
plt.plot(history_rmsprop.history['accuracy'], label='train')
plt.plot(history_rmsprop.history['val_accuracy'], label='test')

plt.title('Model Accuracy: RMSProp optimizer, Relu Activation')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='best')
plt.show()
```



Transfer Learning:

- Prepare the data for the transfer learning algorithm
- Freeze the top layers of the pre-trained model
- Add a dense layer at the end of the pre-trained model followed by a dropout layer
- Add the final output layer with the SoftMax activation function
- Take the loss function as categorical cross-entropy
- Take Adam as an optimizer
- Use early stopping with the patience of two epochs and monitor the validation loss which is set as minimum mode
- Try with fifteen number epochs
- Train the model using the generator and test the accuracy of the test data at every epoch
- Plot the training and validation accuracy, and the loss
- Observe the precision, recall the F1-score for all classes for both grayscale and color models, and determine if the model's classes are good

```

NUM_CLASSES = 7

model_tf = Sequential()
model_tf.add(resnet_v2.ResNet50V2(include_top=False, pooling='max', weights='imagenet'))
model_tf.add(Dense(NUM_CLASSES, activation='softmax'))

model_tf.layers[0].trainable = False

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/resnet/resnet50v2_weights_tf_dim_ordering_tf_kernels_notop.h5
94668760/94668760 [=====] - 3s 0us/step

# Compile the transfer learning model
adam = keras.optimizers.adam_v2.Adam(learning_rate=1e-3, decay=1e-6)

model_tf.compile(optimizer=adam,
                  loss=keras.losses.categorical_crossentropy,
                  metrics=['accuracy'])

model_tf.summary()

Model: "sequential_3"

```

Layer (type)	Output Shape	Param #
resnet50v2 (Functional)	(None, 2048)	23564800
dense_4 (Dense)	(None, 7)	14343

```

=====
Total params: 23,579,143
Trainable params: 14,343
Non-trainable params: 23,564,800
=====

```

```

early_stopping =
EarlyStopping(patience=2, monitor='val_loss', restore_best_weights=True)

history_tf = model_tf.fit(train_generator,
                           epochs=15,
                           callbacks=early_stopping,
                           batch_size=512,
                           verbose=1,
                           validation_data=validation_generator)

Epoch 1/15
451/451 [=====] - 166s 359ms/step - loss:
3.3520 - accuracy: 0.2992 - val_loss: 2.5123 - val_accuracy: 0.3613
Epoch 2/15

```

```

451/451 [=====] - 159s 353ms/step - loss:
2.3808 - accuracy: 0.3475 - val_loss: 2.3007 - val_accuracy: 0.3956
Epoch 3/15
451/451 [=====] - 155s 344ms/step - loss:
2.0713 - accuracy: 0.3700 - val_loss: 1.7990 - val_accuracy: 0.4042
Epoch 4/15
451/451 [=====] - 156s 346ms/step - loss:
1.9403 - accuracy: 0.3825 - val_loss: 1.6937 - val_accuracy: 0.4179
Epoch 5/15
451/451 [=====] - 156s 345ms/step - loss:
1.8904 - accuracy: 0.3829 - val_loss: 1.7142 - val_accuracy: 0.4563
Epoch 6/15
451/451 [=====] - 156s 346ms/step - loss:
1.8997 - accuracy: 0.3828 - val_loss: 1.5949 - val_accuracy: 0.4547
Epoch 7/15
451/451 [=====] - 154s 342ms/step - loss:
1.8693 - accuracy: 0.3874 - val_loss: 1.6574 - val_accuracy: 0.4331
Epoch 8/15
451/451 [=====] - 156s 345ms/step - loss:
1.8657 - accuracy: 0.3872 - val_loss: 1.6431 - val_accuracy: 0.4570

```

Evaluate the model

```

test_loss, test_acc =
model_tf.evaluate(validation_generator, verbose=1)

```

```

print('Model Accuracy', test_acc)
print('Model Loss', test_loss)

```

```

114/114 [=====] - 11s 92ms/step - loss:
1.5949 - accuracy: 0.4547
Model Accuracy 0.45471271872520447
Model Loss 1.5948704481124878

```

Plot the loss function for the model

```


plt.plot(history_tf.history['loss'], label='train')
plt.plot(history_tf.history['val_loss'], label='test')

```

```


plt.title('Model loss: Transfer Learning')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='best')
plt.show()

```

```
# Plot the accuracy function for the model
plt.plot(history_tf.history['accuracy'], label='train')
plt.plot(history_tf.history['val_accuracy'], label='test')

plt.title('Model Accuracy: Transfer Learning')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='best')
plt.show()
```



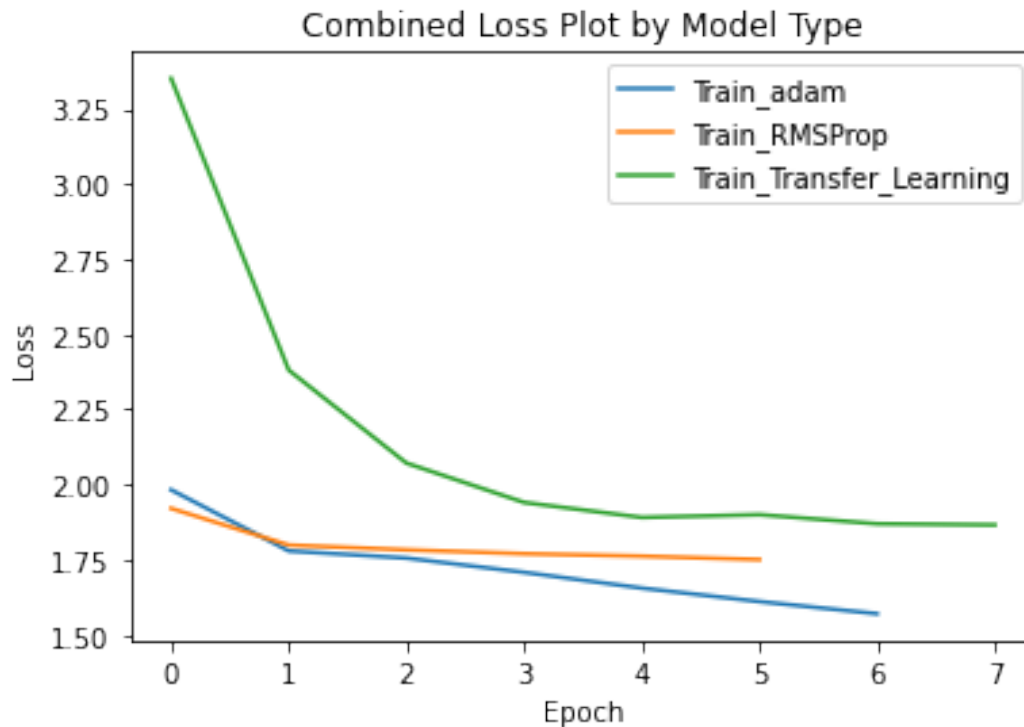
Combined Loss Plot

Plot all loss results on one graph

```
plt.plot(history_adam.history['loss'])
plt.plot(history_rmsprop.history['loss'])
plt.plot(history_tf.history['loss'])

plt.title('Combined Loss Plot by Model Type')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train_adam', 'Train_RMSProp', 'Train_Transfer_Learning'],
           loc='best')

plt.show()
```




Combine Accuracy Plot

```
# Plot all accuracy results on one graph
import matplotlib.pyplot as plt

plt.plot(history_adam.history['accuracy'])
plt.plot(history_rmsprop.history['accuracy'])
plt.plot(history_tf.history['accuracy'])

plt.title('Combined Accuracy Plot by Model Type')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train_adam', 'Train_RMSProp', 'Train_Transfer_Learning'],
           loc='best')
plt.show()
```



Compare all the models on the basis of accuracy, precision, recall, and f1-score


```
y_pred = model_adam.predict(validation_generator)
y_pred = np.argmax(y_pred, axis=1)
print('Confusion Matrix')
print(confusion_matrix(validation_generator.classes, y_pred))
cm = confusion_matrix(validation_generator.classes, y_pred)
plt.figure(figsize=(10,10))
sns.heatmap(cm, annot=True)
plt.xlabel('Predicted')
plt.ylabel('Truth')
```

114/114 [=====] - 4s 38ms/step

Confusion Matrix

```
[[ 33   0  79 344 157 208 139]
 [  2   0   6  37  18  32  16]
 [ 39   0  71 349 184 239 136]
 [ 77   0 130 632 287 406 293]
 [ 51   0  99 428 188 261 189]
 [ 36   0  98 400 183 266 156]
 [ 29   0  69 289 120 176 114]]
```

Text(69.0, 0.5, 'Truth')



```
print(classification_report(validation_generator.classes,y_pred))
```

	precision	recall	f1-score	support
0	0.12	0.03	0.05	960
1	0.00	0.00	0.00	111
2	0.13	0.07	0.09	1018
3	0.25	0.35	0.29	1825
4	0.17	0.15	0.16	1216
5	0.17	0.23	0.20	1139
6	0.11	0.14	0.12	797

accuracy			0.18	7066
macro avg	0.14	0.14	0.13	7066
weighted avg	0.17	0.18	0.17	7066

```

/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))

```

```

y_pred = model_rmsprop.predict(validation_generator)
y_pred = np.argmax(y_pred, axis=1)
print('Confusion Matrix')
print(confusion_matrix(validation_generator.classes, y_pred))
cm = confusion_matrix(validation_generator.classes, y_pred)
plt.figure(figsize=(10,10))
sns.heatmap(cm, annot=True)
plt.xlabel('Predicted')
plt.ylabel('Truth')

```

114/114 [=====] - 5s 40ms/step


Confusion Matrix

```

[[ 27   0  19  556  100   54  204]
 [  6   0   3   60   11    4   27]
 [ 32   0  20  586   85   52  243]
 [ 56   0  31 1036  170  101  431]
 [ 33   0  17  699  140   69  258]
 [ 28   0  22  625  122   78  264]
 [ 24   0  14  465   81   50  163]]

```

Text(69.0, 0.5, 'Truth')



```
print(classification_report(validation_generator.classes,y_pred))
```

	precision	recall	f1-score	support
0	0.13	0.03	0.05	960
1	0.00	0.00	0.00	111
2	0.16	0.02	0.03	1018
3	0.26	0.57	0.35	1825
4	0.20	0.12	0.15	1216
5	0.19	0.07	0.10	1139
6	0.10	0.20	0.14	797

accuracy			0.21	7066
macro avg	0.15	0.14	0.12	7066
weighted avg	0.18	0.21	0.16	7066

```

/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))

```

```

y_pred = model_tf.predict(validation_generator)
y_pred = np.argmax(y_pred, axis=1)
print('Confusion Matrix')
print(confusion_matrix(validation_generator.classes, y_pred))
cm = confusion_matrix(validation_generator.classes, y_pred)
plt.figure(figsize=(10,10))
sns.heatmap(cm, annot=True)
plt.xlabel('Predicted')
plt.ylabel('Truth')

```

114/114 [=====] - 11s 94ms/step

Confusion Matrix

```

[[ 60   8  77 301 319  85 110]
 [  7   0   5  26  49  11  13]
 [ 66  14  67 323 314  96 138]
 [121  15 145 568 595 172 209]
 [ 74   3  82 372 398 116 171]
 [ 57   5  86 353 390  92 156]
 [ 51   7  66 250 263  70  90]]

```

Text(69.0, 0.5, 'Truth')



```
print(classification_report(validation_generator.classes,y_pred))
```

	precision	recall	f1-score	support
0	0.14	0.06	0.09	960
1	0.00	0.00	0.00	111
2	0.13	0.07	0.09	1018
3	0.26	0.31	0.28	1825
4	0.17	0.33	0.22	1216
5	0.14	0.08	0.10	1139
6	0.10	0.11	0.11	797

accuracy			0.18	7066
macro avg	0.13	0.14	0.13	7066
weighted avg	0.17	0.18	0.16	7066

1. CNN Architecture with Adam optimizers Accuracy: 41.80%
 2. CNN customized architecture with RMSProp Optimizers Accuracy: 31.47%
 3. Transfer Learning with Resnet50v2 pretrained model with Adam optimizer Accuracy: 45.47%
- The Transfer Learning with Resnet50v2 performs well as compared to CNN architecture with Adam optimizer and CNN customized architecture with RMSProp optimizer.